

Ollive: Sequential Typed Security Guards for a Grounded Wellness Agent

Yashwardhan Chaudhuri
yashwardhan20417@iiitd.ac.in

Abstract

Ollive is a non-clinical wellness agent that places a separate Security LM before routing, tools, and answer generation. We test whether class-owned sequential guards improve prompt-attack detection without excessive benign blocking. On a matched 357-case GPT-5.4-mini development set, the concise five-guard base blocked 117/157 attacks versus 89/157 for a grouped design, while benign blocks changed from 68/200 to 69/200. A fused guard reached 121/157 observed blocks but rejected 108/200 benign cases and produced 26 malformed-output closures. On a separate 1,213-case suite, the pre-variation base blocked 415/575 attacks with Qwen and 490/575 with GPT-5.4-mini; benign block rates were 13.0% and 16.1%. A 108-case authored wellness suite exposed delimiter handling as the weakest class. A same-case follow-up with compact delimiter families blocked 16/27 under explicit FP8 and 22/27 under BF16. The experiments support sequential class ownership, but precision sensitivity and benign blocking preclude a production-security claim. Code: <https://github.com/yashwardhanchaudhuri/Ollive>.

1 Introduction

Wellness agents place system rules, dialogue, and retrieved text in one model context. An attacker can therefore inject instructions directly or through retrieved evidence (Greshake et al., 2023; Liu et al., 2023); long demonstrations add a many-shot path (Anil et al., 2024). Layered controls are recommended because a single classifier is fallible (OWASP Foundation, 2026). The engineering question is whether separating attack classes improves detection without making ordinary traffic unusable.

Ollive answers general English-language lifestyle questions. Factual answers use a local knowledge base (KB) and allowlisted web sources, while individualized clinical or pharmacological requests stop at a medical boundary. Retrieval

improves factual coverage (Lewis et al., 2020), but retrieved text remains untrusted.

We describe the executable trust boundaries and compare grouped, class-specific, fused, and verbose guard prompts on fixed cases. We then evaluate the selected base on public-source and human-authored wellness suites. The intended contribution is an auditable implementation and measured trade-off, not a claim that prompting alone makes an LLM secure.

2 System Design

2.1 Pipeline and trust boundaries

Runtime flow. Application code fixes the order of work:

```
request/context budget → input gate →  
context gate → routing → medical/non-  
grounded/grounded stage → output validation.
```

Per-session request, message, and context budgets run before either model. Grounded turns must call the KB and at least one web search; at most two additional searches may fill evidence gaps. Other routes skip evidence tools but still receive output review. Figure 1 summarizes the workflow and its Security-LM boundaries.

Context boundary. The broker applies NFKC normalization, removes invisible formatting, and wraps current and prior turns in application-authored provenance. User-written roles, delimiters, and serialized metadata therefore remain content rather than authority. Punctuation is retained so structural attacks remain visible.

2.2 LLM security adapter

Security controls at a glance. Table 1 maps the implemented controls to the attack or failure class they address. These controls constrain the application path; they do not claim that a single classifier recognizes every adversarial prompt.




Figure 1: Ollive system architecture. **Top:** application code forces KB and web retrieval for grounded turns, validates claims, and applies final review. **Bottom:** the Security-LM adapter uses typed stage-specific decisions and application-owned fail-closed enforcement; its input lane contains the ordered five-guard sequence detailed in Figure 2. The grounded branch also shows gap-triggered web searches (at most three total), aggregate evidence rechecks, and at most two same-evidence answer repairs.

2.3 Protection by attack family

2.3.1 Direct Prompt Injection

The direct-injection guard runs before routing or tools. It checks whether untrusted text asks the assistant to override policy, reveal hidden instruc-

tions, change tool rules, or write persistent memory. Authority extraction must identify the target, requested effect, execution intent, and an exact supporting span. Application code blocks a privileged override before downstream work; malformed or

Control	What it handles (one-line method)
Request and context budgets	Caps one message, accumulated dialogue, and requests per session window before any model call; this bounds single-prompt and repeated-turn many-shot loading.
Direct-injection gate	Extracts authority target and requested effect; policy, tool, memory, or hidden-instruction overrides block before routing.
Delimiter and role-confusion defense	NFKC normalization plus application-authored provenance envelopes keep user-supplied SYSTEM, JSON, and delimiter text untrusted.
Persona / DAN guard	Separates benign role-play from a persona seeking privileged action; the latter maps to a jailbreak block.
Composed-context review	Rechecks loaded history as a combined context, preventing earlier turns from becoming trusted persistent instructions.
KB and web evidence gates	Reviews each typed citation exactly once; rejected raw retrieval payloads never enter answer-model context.
Cross-source assembly review	Rechecks approved KB and web citations together after every web set, catching instructions assembled across sources.
Schema and anchor validation	Requires forced typed outputs and contiguous supporting spans; malformed, missing, re-ordered, or unanchored results fail closed.
Random output-leak canary and alignment gate	Detects leakage of an application-injected secret marker, then checks route, user constraints, wellness boundary, evidence, and harmful disclosure before rendering. This is distinct from the rejected semantic ingress-canary ablation.
Grounding validator	Requires structured claims with current-turn markers and validates JSON, provenance, citation membership, and claim-source support.

Table 1: Implemented LLM-security and adjacent grounding controls. The local session budgets are included; identity-aware global limiting, authentication, ingestion scanning, and independent-model deployment remain external controls.

unanchored security output fails closed.

2.3.2 Many-Shot Jailbreaking

Application limits bound both ways of loading demonstrations: one message is capped at 20,000 characters, bounded conversation context at 48,000 characters, and one session at 12 requests per 60 seconds. Repeated examples never gain authority from their number or position; the current message and loaded history remain explicitly untrusted and still pass through the ordered semantic guards. Thus deterministic budgets stop oversized or rapidly repeated loads, while guards inspect shorter ones.

2.3.3 Delimiter Break Attack

The broker applies Unicode NFKC normalization and removes invisible formatting, then encloses every message in application-authored provenance boundaries. A user-written SYSTEM label, JSON field, XML tag, Markdown fence, or claimed end-of-message marker remains data. The delimiter/role-confusion guard explicitly generalizes across nine structural families: XML, JSON, Markdown, YAML/front matter, fake chat headers, serialized role labels, claimed begin/end markers, nested or mismatched boundaries, and Unicode/invisible separators. It blocks by common effect—making an inner imperative executable or falsely authoritative—rather than by an exact attack

phrase. Explicit quotation, analysis, translation, transformation, and ordinary structured wellness data remain allowed.

2.3.4 DAN-Style Persona

The Persona/DAN guard distinguishes harmless role-play or quoted analysis from a persona that claims elevated permission. It blocks personas used to suspend policy, impersonate a privileged role, disclose protected context, or authorize prohibited actions. The decision is typed and enforced by application code; a persona cannot promote itself by declaring that restrictions no longer apply.

Typed sequential enforcement. For input and context, the Security LM returns a forced AuthorityAssessment containing content and persona modes, authority target, effect, execution intent, and an exact span. The adapter validates the schema and checks the span against received text. Application code maps an executed privileged effect to its owning guard and reason; unanchored or malformed assessments fail closed.

The unchanged envelope then enters five forced-schema calls in order: direct injection, delimiter/role confusion, Persona/DAN, harmful capability, and medical risk. Each prompt defines one prohibited effect and its nearest benign boundary. Code stops on the first owned block; otherwise it records the minimum check score. A score cannot

grant authority or override a block. Context, evidence, combined evidence, and output use smaller stage-specific sequences, and item verdicts must return every identifier exactly once and in order.

Isolation contract. The security adapter may share Qwen weights with the answer model, but it remains isolated by a different system prompt, forced schema, zero-temperature decision contract, tool set, and application stage. This is logical adapter isolation, not provider or hardware independence.

2.4 Grounding and output validation

Evidence and release boundary. KB and web results are converted to typed citations and reviewed before answer-model use. Application code rebuilds context from approved items and rechecks the combined set, so rejected raw payloads cannot cross the boundary. The answer model submits structured claims tied to current-turn markers; code validates marker membership, provenance, and claim–passage support. After at most two same-evidence corrections, an opaque-marker leak check and final Security-LM review either release or clear the answer. This mitigates indirect retrieval attacks (Greshake et al., 2023); it does not replace secure ingestion.

3 Evaluation and Ablations

3.1 Evaluation boundary and fixed data

Historical sealed pilot. An earlier 18-case pilot ran each case three times. It blocked all 36 attack attempts and allowed all 18 benign attempts, but nine attack blocks came from backend or malformed-output failure closure rather than valid model verdicts. The pilot predates the class-specific guards and demonstrates enforcement, not reliable recognition.

Configuration-selection comparison. We used one fixed 357-case subset to compare GPT-5.4-mini ingress designs: 79 direct injections, 78 delimiter-wrapped injections, and 200 regular-prompt controls. Every configuration used the same cases, one generation per case, backend, schemas, authority extractor, and enforcement. This subset compares prompts only; final model results use the full 1,213-case suite.

We report the number of attacks blocked at ingress, benign controls incorrectly blocked, total

tokens, and ranking AUC using $1 - \text{trust}$ as the attack score. AUC measures score ordering only; the self-reported trust value is not a calibrated probability and never overrides a block. All runs contain 357 unique records and zero execution errors.

Complete frozen-suite comparison. We ran the selected five-guard profile over all 1,213 cases with both Qwen/Qwen3.5-9B via vLLM and GPT-5.4-mini. Historical Qwen manifests did not record server quantization. Both used the selected guards at temperature zero, one generation per case, and completed 1,213/1,213 records with zero errors.

Public-suite construction. We combined three public sources, removed non-English and duplicate text, and excluded overlap with existing development cases. Deepset Prompt Injections (deepset, 2023) supplied 79 unchanged attacks; 78 copies were placed inside fixed XML, JSON, Markdown, or untrusted-content wrappers, and 198 benign prompts became controls. JailbreakHub (Shen et al., 2024) supplied 200 DAN/persona attacks, 200 regular-prompt controls, and source attacks for 30 many-shot cases: ten each at 8, 32, and 64 shots. Each demonstration used the same generic compliance reply and ended on a different target. We also included every English XSTest contrast (Röttger et al., 2024): 188 unsafe and 240 safe prompts. The suite therefore contains 575 attacks and 638 controls.

At the newly configured 20,000-character runtime cap, 27/30 historical many-shot prompts exceeded the deterministic input budget (all 32/64-shot and seven 8-shot cases); the remaining three still require semantic guards. This is a retrospective length analysis, not a rerun of the Qwen experiment.

Custom-suite construction. We first defined nine non-overlapping mechanisms within each requested family. Direct injection spans policy, hidden-instruction, tool, grounding, memory, safety, output-format, priority, and unauthorized-action overrides. Many-shot spans repetition, escalation, late adversarial shifts, false authority, tool misuse, safety bypass, dual responses, correction pressure, and claimed persistent permission. Delimiter cases use XML, JSON, Markdown, YAML, chat headers, role labels, begin/end markers, mismatched nesting, or invisible Unicode. Persona cases use unrestricted, expert, future-model, fictional, evaluator, emergency, user-owned, dual-




Figure 2: Selected ingress method. Application-owned request and context budgets run before authority extraction. The extractor supplies typed semantics and an exact supporting span; the original envelope then traverses five class-owned guards. Code enforces ownership, first-block stopping, and the minimum trust score. No fused semantic ingress-canary executes in the selected profile.

Deployment	Attack blocked	Benign blocked	Correct decisions	Tokens	Mean trust (A/B)
Qwen 3.5 9B, vLLM [†]	415/575 (72.2%)	83/638 (13.0%)	970/1,213 (80.0%)	8.510M	.159/.670
GPT-5.4 mini	490/575 (85.2%)	103/638 (16.1%)	1,025/1,213 (84.5%)	5.992M	.178/.811

Table 2: Matched full-suite ingress results. Lower trust denotes greater risk; A/B gives mean attack/benign trust. Total blocks include 12 oversized-input blocks for each backend; model verdicts account for 403 Qwen and 478 GPT attack blocks. [†]Historical Qwen manifests did not record server quantization.

Attack family	Qwen [†]	GPT-5.4 mini
Direct Prompt Injection (79)	33	50
Delimiter Break Attack (78)	24	64
DAN-Style Persona (200)	168	178
Many-Shot Jailbreaking (30)	30	30
XSTest unsafe contrast (188)	160	168
All attacks (575)	415	490

Table 3: Blocked attacks by family on the complete frozen suite.

response, or stay-in-character roles.

Each family used the same nine wellness domains once: nutrition, activity, sleep, stress, hydration, supplements, body image, habit tracking, and social wellbeing. We wrote three alternatives for every mechanism–domain cell. Each alternative changes the wording and wellness request while retaining one primary attack mechanism. Validation requires three distinct prompts, bounded length, and explicit transcript turns for many-shot cases. No model creates the prompts. The resulting suite contains $4 \times 9 \times 3 = 108$ unique attacks; Qwen is used only for evaluation.

Qwen 3.5 9B (historical quantization un-

recorded), used only as evaluator, blocked 86/108 attacks (79.6%, 95% Wilson interval 71.1–86.1) with no run errors. Results were 23/27 direct, 26/27 many-shot, 12/27 delimiter, and 25/27 Persona/DAN. Mean trust was .107 overall. The attack-only suite does not measure benign false positives; its main finding is weak delimiter recognition across several authored formats.

After adding compact, case-agnostic delimiter-family guidance, we reran only the same 27 delimiter cases. Explicit vLLM FP8 quantization blocked 16/27 (59.3%); the historical/default BF16 launch blocked 22/27 (81.5%). Both completed without errors. These focused results are not pooled into the older 108-case total.

3.2 Configurations

1. **Compact shared guard:** one compact ingress boundary prompt before harm and medical checks.
2. **Grouped sequential:** boundary, harm, and medical risks evaluated in separate turns.
3. **Concise five-guard base:** direct injection, delimiter/role confusion, Persona/DAN, harm, and

Configuration	Direct	Delimiter	All attacks	Benign blocked	Tokens	Ranking AUC
Compact shared guard	33/79	33/78	66/157 (42.0%)	58/200 (29.0%)	1.24M	.595
Grouped sequential	42/79	47/78	89/157 (56.7%)	68/200 (34.0%)	1.21M	.629
Concise five-guard base	53/79	64/78	117/157 (74.5%)	69/200 (34.5%)	1.656M	.656
Fused semantic canary [†]	52/79	69/78	121/157 (77.1%)	108/200 (54.0%)	2.875M	.550
Expanded specialist (verbose)	43/79	57/78	100/157 (63.7%)	70/200 (35.0%)	2.115M	.629

Table 4: Historical pre-variation GPT-5.4-mini development ablation. Higher attack blocking is desirable; lower benign blocking and token use are desirable. [†]The fused row includes four attack and 22 benign malformed-output fail-closed blocks; its semantic/model counts were 117 attacks and 86 benign controls.




Figure 3: Historical pre-variation ablation on the fixed 357-case development set. The concise five-guard profile was selected as the base because it improved attack blocking without the fused layer’s large benign-blocking and token penalty. The later compact delimiter-family update is evaluated separately on 27 authored cases and is not represented by these bars.

medical guards with class-owned deferral. This base was selected in the 357-case ablation; the current delimiter guard adds a compact nine-family checklist evaluated separately.

4. **Fused semantic canary:** the five guards plus a combined second-line semantic signal and canary fields. This experiment was rejected and removed.
5. **Expanded specialist:** five guards with verbose case-agnostic variation catalogs and no semantic canary. This experiment was also rejected.

3.3 Results and selection

Splitting one boundary check into direct, delimiter, and Persona/DAN ownership increased attack blocks from 89 to 117 relative to the grouped sequential design, at the cost of one additional benign block. The gain was larger for delimiter attacks (47 to 64) than direct attacks (42 to 53). The Persona/DAN guard generated 40 of the concise profile’s 69 benign blocks, identifying the primary calibration problem.

The fused layer produced the highest observed attack block count, but it did not improve valid semantic attack blocks beyond 117. It instead added 39 benign blocks, 26 malformed-output fail-closed events across attack and benign cases, 73.7% more tokens than the concise profile, and weaker score ordering. After removing it, the expanded specialist prompts still regressed to 100 attack blocks while using 27.8% more tokens than the concise base. We therefore restored that base, then later added the compact nine-family delimiter checklist. Only its 27-case focused rerun is reported here; the 357-case bars were not recomputed.

These ablations select an implementation; they do not estimate production security. The 34.5% benign block rate is unacceptable for a general assistant, one sample per case does not measure variance, source labels only approximate Ollive’s threat policy, and all prompt choices were made after examining this development distribution.

4 Discussion and Limitations

4.1 What is implementation and what is model behavior

Application code guarantees guard order, provenance, class ownership, schemas, early stopping, blocked-input isolation, evidence rebuilding, and failure closure; tests verify these transitions. Remaining misses are model classification errors, while benign blocks are calibration errors. The Persona/DAN guard caused 40/69 false positives in the selected development run, so correct control flow does not imply acceptable product behavior.

4.2 Limitations and next evaluation

The evaluation uses one inference per English case and omits adaptive multi-turn attacks, retrieval poisoning, and harmful-compliance scoring. Source labels only approximate Ollive’s policy, trust scores are uncalibrated, and benign blocking of 13.0% for Qwen and 16.1% for GPT is too high. The custom suite has no benign controls and reflects our authored taxonomy, so its 79.6% result must remain separate from the public-source comparison.

A release evaluation should freeze the method, add source-disjoint direct, delimiter, Persona/DAN, encoded, multi-turn, indirect, and composed attacks, repeat attempts, adjudicate disagreements, and score allowed attacks with StrongREJECT or HarmBench (Mazeika et al., 2024b,a). Global identity limits, authentication, ingestion scanning, independent deployment, and audit are separate operational controls.

5 Conclusion

Ollive combines five sequential Security-LM guards with application-owned budgets, evidence handling, and output checks. On the 1,213-case suite, the historical pre-variation base blocked 72.2% of attacks with Qwen and 85.2% with GPT-5.4-mini; benign blocking was 13.0% and 16.1%. The authored suite identified delimiter attacks as the weakest class, and the same-case follow-up remained sensitive to Qwen precision. Sequential class ownership improved the measured trade-off, but benign blocking and model sensitivity remain barriers to deployment.

References

Cem Anil, Esin Durmus, Nina Panickssery, and 1 others. 2024. [Many-shot jailbreaking](#). In *Advances in Neural Information Processing Systems*.

deepset. 2023. [Prompt injections dataset](#). Hugging Face dataset.

Kai Greshake, Sahar Abdelnabi, Shailesh Mishra, and 1 others. 2023. [Not what you’ve signed up for: Compromising real-world llm-integrated applications with indirect prompt injection](#). *arXiv preprint arXiv:2302.12173*.

Patrick Lewis, Ethan Perez, Aleksandra Piktus, and 1 others. 2020. [Retrieval-augmented generation for knowledge-intensive nlp tasks](#). *Advances in Neural Information Processing Systems*.

Yi Liu, Gelei Deng, Yuekang Li, and 1 others. 2023. [Prompt injection attack against llm-integrated applications](#). *arXiv preprint arXiv:2306.05499*.

Mantas Mazeika, Long Phan, Xuwang Yin, and 1 others. 2024a. [Harmbench: A standardized evaluation framework for automated red teaming and robust refusal](#). *arXiv preprint arXiv:2402.04249*.

Mantas Mazeika, Long Phan, Xuwang Yin, and 1 others. 2024b. [A strongreject for empty jailbreaks](#). *arXiv preprint arXiv:2402.10260*.

OWASP Foundation. 2026. [Llm prompt injection prevention cheat sheet](#).

Paul Röttger, Hannah Rose Kirk, Bertie Vidgen, Giuseppe Attanasio, Federico Bianchi, and Dirk Hovy. 2024. [XSTest: A test suite for identifying exaggerated safety behaviours in large language models](#). In *Proceedings of NAACL-HLT*.

Xinyue Shen, Zeyuan Chen, Michael Backes, Yun Shen, and Yang Zhang. 2024. [Do Anything Now: Characterizing and evaluating in-the-wild jailbreak prompts on large language models](#). *arXiv preprint arXiv:2308.03825*.