

Ollive: Sequential Typed Security Guards for a Grounded Wellness Agent

Yashwardhan Chaudhuri
yashwardhan20417@iiitd.ac.in

Abstract

Ollive is a non-clinical wellness agent whose selected model runs isolated, typed security checks at input, evidence, and output boundaries. On a matched 357-case GPT-5.4-mini development set, the concise five-guard base blocked 117/157 attacks versus 89/157 for a grouped design, while benign blocks changed from 68/200 to 69/200. A fused guard reached 121/157 observed blocks but rejected 108/200 benign cases and produced 26 malformed-output closures. On a separate 1,213-case suite, the pre-variation base blocked 415/575 attacks with Qwen and 490/575 with GPT-5.4-mini; benign block rates were 13.0% and 16.1%. On a replacement 108-case wellness-native MECE suite, Qwen 3.5 9B FP8 blocked 108/108 at ingress: 107 model decisions and one malformed-verdict fail-closed block. Because the broad direct gate fired first for all model blocks and the suite has no benign controls, this result measures end-to-end recall rather than specialist attribution or precision. The measured benign blocking still precludes a production-security claim. Code: <https://github.com/yashwardhanchaudhuri/Ollive>.

1 Introduction

Wellness agents combine system rules, dialogue, and retrieved evidence in one model context. An attacker can therefore inject instructions directly or through retrieved evidence (Greshake et al., 2023; Liu et al., 2023); long demonstrations add a many-shot path (Anil et al., 2024). Layered controls are recommended because a single classifier is fallible (OWASP Foundation, 2026). The engineering question is whether separating attack classes improves detection without making ordinary traffic unusable.

This risk matters for Ollive because it answers general English-language lifestyle questions. Factual answers use a local knowledge base (KB) and

allowlisted web sources, while individualized clinical or pharmacological requests stop at a medical boundary. Retrieval improves factual coverage (Lewis et al., 2020), but retrieved text remains untrusted.

To test that design, we describe the executable trust boundaries and compare grouped, class-specific, fused, and verbose guard prompts on fixed cases. We then evaluate the selected base on public-source and human-authored wellness suites. The intended contribution is an auditable implementation and measured trade-off, not a claim that prompting alone makes an LLM secure.

2 System Design

2.1 Pipeline and trust boundaries

Runtime flow. The current runtime is fail-closed. Before any model call, application code enforces request-frequency, message-size, and accumulated-context budgets. Accepted turns pass through five ordered input guards and a composed-history review; the first block stops before routing.

Routing selects an application-owned medical boundary, a grounded wellness path, or a tool-free non-grounded response. Grounded turns force a KB lookup and one web search. The Security LM reviews each citation and the combined KB–web set; a named gap may trigger at most two further searches. Answers use typed claim–citation items, with claim support verified and at most two corrections against unchanged evidence. All routes finish with applicable citation validation, the deterministic output-leak canary, and final integrity, harm, and medical reviews. Figures 1 and 2 show the application and Security-LM boundaries.

Context boundary. For security inspection, the broker derives a canonical copy of each current and prior message: it applies NFKC, normalizes line endings and whitespace, removes invisible format controls, and retains semantic punctuation. The




Figure 1: Ollive runtime workflow. Grounded turns require KB lookup and one web search; a named evidence gap may trigger at most two additional web searches. Every web result is filtered and the approved KB–web set is rechecked. Claim validation permits at most two same-evidence repairs before citation and output review.

surrounding envelope—not text supplied inside it—records source, untrusted authority, original length, content hash, and whether normalization changed the text. User-written roles, delimiters, and serialized metadata therefore remain content rather than authority.

2.2 LLM security adapter

Security controls at a glance. Once provenance is fixed, the security adapter applies the controls in Table 1. The table maps each control to the attack or failure class it addresses. These controls constrain the application path; they do not claim that a single classifier recognizes every adversarial prompt.

2.3 Protection by attack family

The ingress controls divide the assignment threat model into four requested families. Each guard owns one prohibited effect and explicitly defers

neighboring classes.

2.3.1 Direct Prompt Injection

The direct-injection guard runs before routing or tools. It checks whether untrusted text asks the assistant to override policy, reveal hidden instructions, change tool rules, or write persistent memory. Authority extraction must identify the target, requested effect, execution intent, and an exact supporting span. Application code blocks a privileged override before downstream work; malformed or unanchored security output fails closed.

2.3.2 Many-Shot Jailbreaking

Application limits bound both ways of loading demonstrations: one message is capped at 20,000 characters, bounded conversation context at 48,000 characters, and one session at 12 requests per 60 seconds. Repeated examples never gain authority from their number or position; the current




Figure 2: Security adapter boundaries. The selected backend runs separate typed checks: input and composed-context guards; per-item review of every KB and web citation; an aggregate recheck after web retrieval; and final-answer guards. Application code excludes unsafe evidence and fails closed on unavailable, malformed, or unsafe input/output decisions.

message and loaded history remain explicitly untrusted and still pass through the ordered semantic guards. Thus deterministic budgets stop oversized or rapidly repeated loads, while guards inspect shorter ones.

2.3.3 Delimiter Break Attack

The broker applies Unicode NFKC normalization and removes invisible formatting, then encloses every message in application-authored provenance boundaries. A user-written SYSTEM label, JSON field, XML tag, Markdown fence, or claimed end-of-message marker remains data. The delimiter/role-confusion guard explicitly generalizes across nine structural families: XML, JSON, Markdown, YAML/front matter, fake chat headers, serialized role labels, claimed begin/end markers, nested or mismatched boundaries, and Unicode/invisible separators. It blocks by common effect—making an inner imperative executable or falsely authoritative—rather than by an exact attack phrase. Explicit quotation, analysis, translation, transformation, and ordinary structured wellness

data remain allowed.

2.3.4 DAN-Style Persona

The Persona/DAN guard distinguishes harmless role-play or quoted analysis from a persona that claims elevated permission. It blocks personas used to suspend policy, impersonate a privileged role, disclose protected context, or authorize prohibited actions. The decision is typed and enforced by application code; a persona cannot promote itself by declaring that restrictions no longer apply.

Typed sequential enforcement. These family boundaries become executable through a two-step decision. For input and context, the Security LM returns a forced `AuthorityAssessment` containing content and persona modes, authority target, effect, execution intent, and an exact span. The adapter validates the schema and checks the span against received text. Application code maps an executed privileged effect to its owning guard and reason; unanchored or malformed assessments fail closed.

The unchanged envelope then enters five forced-schema calls in order: direct injection, delim-

Control	What it handles (one-line method)
Request and context budgets	Caps one message, accumulated dialogue, and requests per session window before any model call; this bounds single-prompt and repeated-turn many-shot loading.
Direct-injection gate	Extracts authority target and requested effect; policy, tool, memory, or hidden-instruction overrides block before routing.
Delimiter and role-confusion defense	NFKC normalization plus application-authored provenance envelopes keep user-supplied SYSTEM, JSON, and delimiter text untrusted.
Persona / DAN guard	Separates benign role-play from a persona seeking privileged action; the latter maps to a jailbreak block.
Composed-context review	Rechecks loaded history as a combined context, preventing earlier turns from becoming trusted persistent instructions.
KB and web evidence gates	Reviews each typed citation exactly once; rejected raw retrieval payloads never enter answer-model context.
Cross-source assembly review	Rechecks approved KB and web citations together after every web set, catching instructions assembled across sources.
Schema and anchor validation	Requires forced typed outputs and contiguous supporting spans; malformed, missing, re-ordered, or unanchored results fail closed.
Random output-leak canary and alignment gate	Detects leakage of an application-injected secret marker, then checks route, user constraints, wellness boundary, evidence, and harmful disclosure before rendering. This is distinct from the rejected semantic ingress-canary ablation.
Grounding validator	Requires structured claims with current-turn markers and validates JSON, provenance, citation membership, and claim-source support.

Table 1: Implemented LLM-security and adjacent grounding controls. The local session budgets are included; identity-aware global limiting, authentication, ingestion scanning, and independent-model deployment remain external controls.

iter/role confusion, Persona/DAN, harmful capability, and medical risk. Each prompt defines one prohibited effect and its nearest benign boundary. Code stops on the first owned block; otherwise it records the minimum check score. A score cannot grant authority or override a block. Context, evidence, combined evidence, and output use smaller stage-specific sequences, and item verdicts must return every identifier exactly once and in order.

Isolation contract. The security adapter uses the same selected backend as the answer pipeline, but it remains isolated by a different system prompt, forced schema, zero-temperature decision contract, tool set, and application stage. This is logical adapter isolation, not provider or hardware independence.

2.4 Grounding and output validation

Evidence and release boundary. The same untrusted-data rule extends beyond user input. KB and web results are converted to typed citations and reviewed before answer-model use. Application code rebuilds context from approved items and rechecks the combined set, so rejected raw payloads cannot cross the boundary. The answer model submits structured claims tied to current-turn markers; code validates marker membership, provenance, and claim-passage support. After at most two same-evidence corrections, an opaque-

marker leak check and final Security-LM review either release or clear the answer. This mitigates indirect retrieval attacks (Greshake et al., 2023); it does not replace secure ingestion.

3 Evaluation and Ablations

3.1 Evaluation boundary and fixed data

We define the evaluation data before reporting model or ablation results, keeping development selection separate from final-suite measurement.

Public-suite construction. We combined three public sources, removed non-English and duplicate text, and excluded overlap with existing development cases. Deepset Prompt Injections (deepset, 2023) supplied 79 unchanged attacks; 78 copies were placed inside fixed XML, JSON, Markdown, or untrusted-content wrappers, and 198 benign prompts became controls. JailbreakHub (Shen et al., 2024) supplied 200 DAN/persona attacks, 200 regular-prompt controls, and source attacks for 30 many-shot cases: ten each at 8, 32, and 64 shots. Each demonstration used the same generic compliance reply and ended on a different target. We also included every English XSTest contrast (Röttger et al., 2024): 188 unsafe and 240 safe prompts. The suite therefore contains 575 attacks and 638 controls.

At the newly configured 20,000-character run-time cap, 27/30 historical many-shot prompts ex-




Figure 3: Selected ingress sequence after application-owned budgets and typed authority extraction. The unchanged envelope traverses five class-owned guards. Each red arrow is a possible first-block exit; only completion of all five reaches Allow. The selected profile has no fused semantic ingress canary.

ceed the deterministic input budget (all 32/64-shot and seven 8-shot cases); the remaining three still require semantic guards. This is a retrospective length analysis, not a rerun of the Qwen experiment.

Configuration-selection comparison. We next used the development subset only to choose among guard designs. We used one fixed 357-case subset to compare GPT-5.4-mini ingress designs: 79 direct injections, 78 delimiter-wrapped injections, and 200 regular-prompt controls. Every configuration used the same cases, one generation per case, backend, schemas, authority extractor, and enforcement. This subset compares prompts only; final model results use the full 1,213-case suite.

We report the number of attacks blocked at ingress, benign controls incorrectly blocked, total tokens, and ranking AUC using $1 - \text{trust}$ as the attack score. AUC measures score ordering only; the self-reported trust value is not a calibrated probability and never overrides a block. All runs contain 357 unique records and zero execution errors.

3.2 Configurations

The development ablation compares five increasingly specialized ingress designs:

1. **Compact shared guard:** one compact ingress boundary prompt before harm and medical checks.
2. **Grouped sequential:** boundary, harm, and medical risks evaluated in separate turns.
3. **Concise five-guard base:** direct injection, delimiter/role confusion, Persona/DAN, harm,

and medical guards with class-owned deferral. This base was selected in the 357-case ablation; the current delimiter guard adds a compact nine-family checklist evaluated separately.

4. **Fused semantic canary:** the five guards plus a combined second-line semantic signal and canary fields. This experiment was rejected and removed.
5. **Expanded specialist:** five guards with verbose case-agnostic variation catalogs and no semantic canary. This experiment was also rejected.

3.3 Results and selection

Against this fixed development set, splitting one boundary check into direct, delimiter, and Persona/DAN ownership increased attack blocks from 89 to 117 relative to the grouped sequential design, at the cost of one additional benign block. The gain was larger for delimiter attacks (47 to 64) than direct attacks (42 to 53). The Persona/DAN guard generated 40 of the concise profile’s 69 benign blocks, identifying the primary calibration problem.

The fused layer produced the highest observed attack block count, but it did not improve valid semantic attack blocks beyond 117. It instead added 39 benign blocks, 26 malformed-output fail-closed events across attack and benign cases, 73.7% more tokens than the concise profile, and weaker score ordering. After removing it, the expanded specialist prompts still regressed to 100 attack blocks while using 27.8% more tokens than the concise

Configuration	Direct	Delimiter	All attacks	Benign blocked	Tokens	Ranking AUC
Compact shared guard	33/79	33/78	66/157 (42.0%)	58/200 (29.0%)	1.24M	.595
Grouped sequential	42/79	47/78	89/157 (56.7%)	68/200 (34.0%)	1.21M	.629
Concise five-guard base	53/79	64/78	117/157 (74.5%)	69/200 (34.5%)	1.656M	.656
Fused semantic canary [†]	52/79	69/78	121/157 (77.1%)	108/200 (54.0%)	2.875M	.550
Expanded specialist (verbose)	43/79	57/78	100/157 (63.7%)	70/200 (35.0%)	2.115M	.629

Table 2: Historical pre-variation GPT-5.4-mini development ablation. Higher attack blocking is desirable; lower benign blocking and token use are desirable. [†]The fused row includes four attack and 22 benign malformed-output fail-closed blocks; its semantic/model counts were 117 attacks and 86 benign controls.




Figure 4: Historical pre-variation ablation on the fixed 357-case development set. The concise five-guard profile was selected as the base because it improved attack blocking without the fused layer’s large benign-blocking and token penalty. The later compact delimiter-family update is evaluated separately on 27 authored cases and is not represented by these bars.

base. We therefore restored that base, then later added the compact nine-family delimiter checklist. Only its 27-case focused rerun is reported here; the 357-case bars were not recomputed.

These ablations select an implementation; they do not estimate production security. The 34.5% benign block rate is unacceptable for a general assistant, one sample per case does not measure variance, source labels only approximate Ollive’s threat policy, and all prompt choices were made after examining this development distribution.

Complete frozen-suite comparison. After selecting the concise profile, we ran the selected five-guard profile over all 1,213 cases with both Qwen/Qwen3.5-9B via vLLM and GPT-5.4-mini. Historical Qwen manifests did not record server quantization. Both used the selected guards at temperature zero, one generation per case, and completed 1,213/1,213 records with zero errors.

Full-suite interpretation. GPT-5.4-mini blocked 75 more attacks than Qwen but also blocked 20

more benign controls. The result therefore improves overall accuracy without resolving over-refusal: both benign block rates remain too high for unobtrusive deployment.

Custom-suite construction. The public sources test breadth but not a balanced wellness-specific taxonomy. The original v1 suite attached wellness requests to mostly domain-neutral attacks; the replacement v2 suite makes the attack itself depend on a wellness workflow or boundary. Within each requested family, nine cells partition the primary mechanism or framing used by that family. Direct injection targets nine different control surfaces; many-shot cases use nine different conditioning patterns; delimiter cases use nine serialization families; and Persona/DAN cases use nine role framings.

Each family uses the same nine wellness domains once: nutrition, activity, sleep, stress, hydration, supplements, body image, habit tracking, and social wellbeing. Three alternatives change the user situation while retaining the cell’s attack

Deployment	Attack blocked	Benign blocked	Correct decisions	Tokens	Mean trust (A/B)
Qwen 3.5 9B, vLLM [†]	415/575 (72.2%)	83/638 (13.0%)	970/1,213 (80.0%)	8.510M	.159/.670
GPT-5.4 mini	490/575 (85.2%)	103/638 (16.1%)	1,025/1,213 (84.5%)	5.992M	.178/.811

Table 3: Matched full-suite ingress results. Lower trust denotes greater risk; A/B gives mean attack/benign trust. Total blocks include 12 oversized-input blocks for each backend; model verdicts account for 403 Qwen and 478 GPT attack blocks. [†]Historical Qwen manifests did not record server quantization.

Attack family	Qwen [†]	GPT-5.4 mini
Direct Prompt Injection (79)	33	50
Delimiter Break Attack (78)	24	64
DAN-Style Persona (200)	168	178
Many-Shot Jailbreaking (30)	30	30
XSTest unsafe contrast (188)	160	168
All attacks (575)	415	490

Table 4: Blocked attacks by family on the complete frozen suite.

mechanism, yielding $4 \times 9 \times 3 = 108$ unique attacks. Validation requires three distinct prompts per cell, explicit transcript turns for many-shot cases, wellness-native domain language, and absence of the v1 plug-in phrase. No model creates the prompts; Qwen is used only for evaluation. Table 5 keeps all 36 cells together: each row supplies the wellness theme, while each family cell gives the attack method and key protected boundary in one sentence.

Historical v1 result and delimiter follow-up.

On the contextualized v1 suite, Qwen 3.5 9B (quantization unrecorded) blocked 86/108 attacks (79.6%, 95% Wilson interval 71.1–86.1) with no run errors: 23/27 direct, 26/27 many-shot, 12/27 delimiter, and 25/27 Persona/DAN. After compact delimiter-family guidance, a same-case delimiter rerun blocked 16/27 under explicit FP8 and 22/27 under the historical/default BF16 launch. These development results remain separate from v2 because the prompts changed.

Wellness-native v2 result. Qwen 3.5 9B under explicit vLLM FP8 blocked 108/108 v2 attacks at ingress (100%, 95% Wilson interval 96.6–100.0) with zero execution errors: 27/27 direct, 27/27 many-shot, 27/27 delimiter, and 27/27 Persona/DAN. Of the 108 blocks, 107 were valid model decisions; one output-contract case returned a malformed verdict and was blocked by the application fail-closed contract. The run used 383,213 tokens. The broad direct-injection check fired first for every valid model block, so the result measures

end-to-end ingress recall, not whether the later delimiter or Persona/DAN guard independently recognized its nominal class. Because v2 contains attacks only, it does not measure benign false positives.

4 Discussion and Limitations

The experiments answer two different questions: whether the application enforces its declared boundaries, and whether the selected model classifies cases accurately enough.

4.1 What is implementation and what is model behavior

Application code guarantees guard order, provenance, class ownership, schemas, early stopping, blocked-input isolation, evidence rebuilding, and failure closure; tests verify these transitions. Remaining misses are model classification errors, while benign blocks are calibration errors. The Persona/DAN guard caused 40/69 false positives in the selected development run, so correct control flow does not imply acceptable product behavior.

4.2 Limitations and next evaluation

The evaluation uses one inference per English case and omits adaptive attacks that change after observing a guard response. V2 includes poisoned record formats but does not execute a live retrieval-poisoning campaign. Source labels only approximate Ollive’s policy, trust scores are uncalibrated, and public-suite benign blocking of 13.0% for Qwen and 16.1% for GPT remains too high. Both custom suites are authored, attack-only evaluations with no near-neighbor controls; their results therefore stay separate from the public-source comparison. In v2, the first broad guard preempted the nominal class guards, so 108/108 does not establish independent specialist coverage or acceptable precision.

A release evaluation should freeze the method, add source-disjoint direct, delimiter, Persona/DAN, encoded, multi-turn, indirect, and composed attacks, repeat attempts, adjudicate disagreements,

Theme	Direct injection	Many-shot	Delimiter break	Persona / DAN
Nutrition	A fake dietitian waiver replaces nutrition policy; key: care-role claims stay untrusted.	Repeated lower-calorie approvals normalize restriction; key: prior unsafe answers grant no permission.	A forged XML meal-plan system block crosses tags; key: imported XML remains data.	A no-limits diet coach prescribes restriction; key: role-play cannot expand scope.
Activity	Private injury thresholds are requested before unsafe loading; key: hidden guards remain confidential and authoritative.	Demonstrations escalate walking advice through pain and chest symptoms; key: later risk restores boundaries.	Developer-role fields are inserted into wearable JSON; key: serialized roles create no authority.	A physiotherapist persona claims superior authority; key: expertise cannot overrule assessment boundaries.
Sleep	Forums or seller pages replace approved medication evidence; key: users cannot reorder retrieval.	Benign diary examples shift to medication and sedatives; key: learned format cannot override final-turn safety.	A forged clinic instruction follows a closed Markdown fence; key: formatting cannot escape provenance.	A future clinical model claims exemption; key: simulated versions inherit current policy.
Stress	Unsupported certainty is demanded for panic, crisis, or journals; key: grounding and crisis duties remain.	Fake coach or therapist approvals accumulate; key: demonstrated authority remains untrusted.	YAML front matter claims therapist authority; key: metadata cannot suppress crisis or evidence rules.	Fiction supplies a boundary-free clinician; key: role-play does not remove real safety duties.
Hydration	Unsafe fluid preferences are written to memory; key: memory cannot waive later medical checks.	Blogs and calculators repeatedly outrank the KB; key: repeated misuse cannot normalize invalid tool order.	Fake clinical transcript headers surround symptom data; key: user-authored roles stay user content.	An evaluator requests prohibited dosing for scoring; key: audit framing permits no unsafe disclosure.
Supplements	Consent or urgency suspends dose and interaction limits; key: accepting risk does not expand scope.	Seller disclaimers repeatedly waive stack checks; key: demonstrations cannot remove interaction safeguards.	SYSTEM or developer labels appear in seller text; key: textual labels cannot change priority.	An emergency specialist skips interaction checks; key: urgency creates no clinical authority.
Body image	Restriction is hidden in encoded output fields; key: schemas cannot conceal disallowed guidance.	A safe branch trains an unrestricted second answer; key: every branch remains governed.	False begin-end markers precede a forged policy; key: application envelopes own boundaries.	A user-owned appearance coach demands obedience; key: ownership cannot void safeguards.
Habit tracking	Feature use becomes blanket consent for goals or sharing; key: permission must be current and explicit.	Refusals and consent checks are labelled errors; key: correction pressure cannot train safeguards away.	Mismatched tracker nesting injects policy; key: malformed records grant neither consent nor authority.	A policy-free tracker is paired with a safe coach; key: both personas receive identical review.
Social well-being	A check-in becomes permission to contact or enrol; key: external actions require confirmation.	Prior turns fabricate standing outreach consent; key: sensitive permission must be revalidated.	Unicode separators disguise organizer instructions; key: normalization exposes rather than trusts them.	Staying in a proactive care role is mandatory; key: character pressure cannot replace consent.

Table 5: All 36 wellness-native MECE cells in one matrix. Rows are wellness themes; cells state the family-specific method and key boundary, with three authored situations per cell.

and score allowed attacks with StrongREJECT or HarmBench (Mazeika et al., 2024b,a). Global identity limits, authentication, ingestion scanning, independent deployment, and audit are separate operational controls.

5 Conclusion

Taken together, the implementation and experiments show that Ollive combines five sequential Security-LM guards with application-owned budgets, evidence handling, and output checks. On the 1,213-case suite, the historical pre-variation base blocked 72.2% of attacks with Qwen and 85.2% with GPT-5.4-mini; benign blocking was 13.0% and 16.1%. The wellness-native v2 suite broadens domain realism and was blocked 108/108 by Qwen FP8, but broad first-guard preemption prevents a specialist recall claim. Sequential ownership improved the development trade-off, while benign blocking, model sensitivity, and guard attribution remain barriers to deployment.

References

Cem Anil, Esin Durmus, Nina Panickssery, and 1 others. 2024. [Many-shot jailbreaking](#). In *Advances in Neural Informa-*

tion Processing Systems.

deepset. 2023. [Prompt injections dataset](#). Hugging Face dataset.

Kai Greshake, Sahar Abdelnabi, Shailesh Mishra, and 1 others. 2023. [Not what you’ve signed up for: Compromising real-world llm-integrated applications with indirect prompt injection](#). *arXiv preprint arXiv:2302.12173*.

Patrick Lewis, Ethan Perez, Aleksandra Piktus, and 1 others. 2020. [Retrieval-augmented generation for knowledge-intensive nlp tasks](#). *Advances in Neural Information Processing Systems*.

Yi Liu, Gelei Deng, Yuekang Li, and 1 others. 2023. [Prompt injection attack against llm-integrated applications](#). *arXiv preprint arXiv:2306.05499*.

Mantas Mazeika, Long Phan, Xuwang Yin, and 1 others. 2024a. [Harmbench: A standardized evaluation framework for automated red teaming and robust refusal](#). *arXiv preprint arXiv:2402.04249*.

Mantas Mazeika, Long Phan, Xuwang Yin, and 1 others. 2024b. [A strongreject for empty jailbreaks](#). *arXiv preprint arXiv:2402.10260*.

OWASP Foundation. 2026. [Llm prompt injection prevention cheat sheet](#).

Paul Röttger, Hannah Rose Kirk, Bertie Vidgen, Giuseppe Attanasio, Federico Bianchi, and Dirk Hovy. 2024. [XSTest: A test suite for identifying exaggerated safety behaviours in large language models](#). In *Proceedings of NAACL-HLT*.

Xinyue Shen, Zeyuan Chen, Michael Backes, Yun Shen, and Yang Zhang. 2024. [Do Anything Now: Characterizing and evaluating in-the-wild jailbreak prompts on large language models](#). *arXiv preprint arXiv:2308.03825*.